

# Keras

Keras for Deep Learning Research

Feedback is greatly appreciated!

# Unsupervised Learning

- Most of the world's data is unlabeled!
- New articles, movie reviews, user Netflix ratings, images on the internet, etc.
- What approaches can we take to make use of this type of unlabeled data?

# Unsupervised Learning


- Two main approaches
  - Dimensionality Reduction
  - Clustering

# Autoencoder

- The autoencoder is actually a very simple neural network and will feel similar to a multi-layer perceptron model.
- It is designed to reproduce its input at the output layer.
- The key difference between an autoencoder and a typical MLP network is that the number of input neurons is equal to the number of output neurons.


# Autoencoder

- An autoencoder is a neural network that has three layers: an input layer, a hidden (encoding) layer, and a decoding layer.
- The network is trained to reconstruct its inputs, which forces the hidden layer to try to learn good representations of the inputs.




# Autoencoder

- The general structure of an autoencoder, mapping an **input  $x$**  to an **output  $r$**  (called **reconstruction**) through an internal representation or **code  $h$** .
- The autoencoder has two components: the encoder  $f$  (mapping  $x$  to  $h$ ) and the decoder  $g$  (mapping  $h$  to  $r$ ).



# Example Autoencoder




Source: <https://www.udemy.com/deeplearning/learn/v4/overview>



# Autoencoder

- Feed forward network trained to reproduce its input at the output layer.
- Output size is the same as the input layer.




# Autoencoder

**Decoder**

$$\hat{x} = \text{sigm}(c + W^{\text{out}} h(x))$$


**Encoder**

$$h(x) = \text{sigm}(b + Wx)$$



# Autoencoder


- The hidden/internal representation maintains all the information of the input.
- We can use the hidden layer to extract meaningful features.



# Applications of Autoencoder

- Dimensionality Reduction
- Data Denoising
- Image Reconstruction
- Information Retrieval


# Dimensionality Reduction




# Dimensionality Reduction

- Dimensionality reduction allows us to get a lower dimension representation of our data.
- The encoder creates new (fewer) features from the input features.
- For example we can input a 3 dimensional data set and output a 2 dimensional representation of it.

# Dimensionality Reduction



# Data Denoising






# Image Reconstruction

Input



Reconstructed input



Removes the dark cross-bar in the image


# Information Retrieval

- This is related to text domain
- We can use an autoencoder to find low dimensional codes for documents that allow fast and accurate retrieval of similar documents from a large set.

# Autoencoder vs PCA

- PCA is restricted to a linear map, while auto encoders can have nonlinear encoder/decoders.
- A single layer auto encoder with linear transfer function is nearly equivalent to PCA, where nearly means that the  $\mathbf{W}$  found by AE and PCA won't be the same--but the subspace spanned by the respective  $\mathbf{W}$ 's will.

# Stacked Autoencoder





Source: <https://www.udemy.com/deeplearning/learn/v4/overview>

# Stacked Autoencoder

- A stacked autoencoder is a neural network consisting of multiple layers of sparse autoencoders in which the outputs of each layer is wired to the inputs of the successive layer.
- Then the encoding step for the stacked autoencoder is given by running the encoding step of each layer in forward order:  $z^{(l+1)} = W^{(l,1)} a^l + b^{(l,1)}$
- The decoding step is given by running the decoding stack of each autoencoder in reverse order:  $z^{(n+l+1)} = W^{(n-l,2)} a^{(n+l)} + b^{(n-l,2)}$

# Architecture of Autoencoder



# Architecture of Autoencoders

- Both the encoder and decoder are fully-connected feedforward neural networks
- Code is a single layer of an ANN with the dimensionality of our choice.
- The number of nodes in the code layer (**code size**) is a hyperparameter that we set before training the autoencoder.
- Note that the decoder architecture is the mirror image of the encoder.
- The only requirement is the dimensionality of the input and output needs to be the same. Anything in the middle can be played with.

# Hyperparameters of Autoencoder


- Code size
- Number of layers
- Number of nodes per layer
- Loss function




# Types of Autoencoder

- **Undercomplete Autoencoder:** An autoencoder whose **code** dimension is **less than** the input dimension is called undercomplete. Learning an undercomplete representation forces the autoencoder to capture the most salient features of the training data.
- **Overcomplete Autoencoder:** An autoencoder whose **code** dimension is **greater than** the input dimension is called overcomplete.
- **Regularized Autoencoders:** Rather than limiting the size of the code dimension for the sake of feature learning, we can add a loss function to prevent it memorizing the task and the training data.


# Undercomplete and Overcomplete AE




(a) Shallow undercomplete



(b) Shallow overcomplete



(c) Deep undercomplete




(d) Deep overcomplete

# Types of Regularized Autoencoders

- Denoising Autoencoder (DAE)
- Sparse Autoencoder (SAE)
- Variational Autoencoder (VAE)
- Contractive Autoencoder (CAE)

# Denoising Autoencoder




# Denoising Autoencoder

- There is another way to force the autoencoder to learn useful features, which is adding random noise to its inputs and making it recover the original noise-free data.
- This way the autoencoder can't simply copy the input to its output because the input also contains random noise.
- We are asking it to subtract the noise and produce the underlying meaningful data.
- This is called a denoising autoencoder.

# Different kinds of noises

- Salt and Pepper Noise
- Gaussian Noise
- Periodic Noise
- Speckle Noise

# Sparse Autoencoder



Source: <http://www.ericwilkinson.com/blog/2014/11/19/deep-learning-sparse-autoencoders>

# Sparse Autoencoder

- An autoencoder which has a sparsity penalty in the training loss in addition to the reconstruction error.
- We can regularize the autoencoder by using a sparsity constraint such that only a fraction of the nodes would have nonzero values, called active nodes.
- In particular, we add a **penalty term** to the **loss function** such that only a **fraction of the nodes** become **active**.
- This method works even if the code size is large, since only a small subset of the nodes will be active at any time.



# Usage of callbacks

- A callback is a set of functions to be applied at given stages of the training procedure
- You can use callbacks to get a view on internal states and statistics of the model during training

# Usage of callbacks

- **TerminateOnNaN:** Callback that terminates training when a NaN loss is encountered.
- **History:** Callback that records events into a History object
- **ModelCheckpoint:** Save the model after every epoch.
- **EarlyStopping:** Stop training when a monitored quantity has stopped improving.
- **LearningRateScheduler:** Learning rate scheduler.
- **ReduceLROnPlateau:** Reduce learning rate when a metric has stopped improving.

# Use Case: Fashion\_MNIST data set



# Creating model

```
encoding_dim = 32
```

```
# this is our input placeholder
```

```
input_img = Input(shape=(784,))
```

```
# "encoded" is the encoded representation of the input
```

```
encoded = Dense(encoding_dim, activation='relu')(input_img)
```

```
# "decoded" is the lossy reconstruction of the input
```

```
decoded = Dense(784, activation='sigmoid')(encoded)
```

```
# this model maps an input to its reconstruction
```

```
autoencoder = Model(input_img, decoded)
```

```
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy')
```

# Reading data set

- First, we'll configure our model to use a per-pixel binary crossentropy loss, and the Adadelta optimizer:

```
from keras.datasets import fashion_mnist
```

```
import numpy as np
```

```
(x_train, _), (x_test, _) = fashion_mnist.load_data()
```

# Fitting model

- We will normalize all values between 0 and 1 and we will flatten the 28x28 images into vectors of size 784.
- `x_train = x_train.astype('float32') / 255.`  
`x_test = x_test.astype('float32') / 255.`  
`x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))`  
`x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))`

```
autoencoder.fit(x_train, x_train,  
               epochs=5,  
               batch_size=256,  
               shuffle=True,  
               validation_data=(x_test, x_test))
```

```
prediction = autoencoder.predict(X_test[0])
```

# Use case2: Denosing Autoencoder

- The only difference is that we feed a noisy data to the model:
- Noise factor controls the noisiness of images
- clip the values to make sure that the elements of feature vector representing image are between 0 and 1

*#introducing noise*

noise\_factor = 0.5

x\_train\_noisy = x\_train + noise\_factor \* np.random.normal(loc=0.0, scale=1.0, size=x\_train.shape)

x\_test\_noisy = x\_test + noise\_factor \* np.random.normal(loc=0.0, scale=1.0, size=x\_test.shape)

autoencoder.fit(x\_train\_noisy, x\_train,  
 epochs=10,  
 batch\_size=256,  
 shuffle=True,  
 validation\_data=(x\_test\_noisy, x\_test\_noisy))

# Visualizing the train data

```
from matplotlib import pyplot as plt  
plt.imshow(x_train[1].reshape(28,28))  
plt.show()
```

```
from matplotlib import pyplot as plt  
plt.imshow(prediction.reshape(28,28))  
plt.show()
```



# References

- <http://www.wildml.com/2015/12/implementing-a-cnn-for-text-classification-in-tensorflow/>
- <https://towardsdatascience.com/a-walkthrough-of-convolutional-neural-network-7f474f91d7bd>